

Hands-on Lab: Improving Performance of Slow Queries in MySQL



Estimated time needed: 45 minutes

In this lab, you will learn how to improve the performance of your slow queries in MySQL, which can be particularly helpful with large databases.

Objectives

After completing this lab, you will be able to:

1. Use the `EXPLAIN` statement to check the performance of your query
2. Add indexes to improve the performance of your query
3. Apply other best practices such as using the `EXISTS` clause to improve query performance

Software Used in this Lab

In this lab, you will use [MySQL](#). MySQL is a Relational Database Management System (RDBMS) designed to efficiently store, manipulate, and retrieve data.




To complete this lab, you will utilize the MySQL relational database service available as part of the IBM Skills Network Labs (SN Labs) Cloud IDE. SN Labs is a virtual lab environment used in this course.

Database Used in this Lab

The Employees database used in this lab comes from the following source: <https://cse.msu.edu/~cse501/employees/>, under the [CC BY-SA 3.0 license](#).

The following entity relationship diagram (ERD) shows the schema of the Employees database:



The first row of each table is the table name, the rows with keys next to them indicate the primary keys, and the remaining rows are additional attributes.

Exercise 1: Load the Database

Let's begin by retrieving the database and loading it so that it can be used.

1. In the menu bar, select `Terminal` > `New Terminal`. This will open the Terminal.

To download the zip file containing the database, copy and paste the following into the Terminal:

```
1. curl -O https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/datasets/employeesdb.zip
```

```
theia@theiadocker-:/home/project$ wget https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/datasets/employeesdb.zip
--2021-10-12 20:08:23-- https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/datasets/employeesdb.zip
Resolving cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud)... 198.23.119.2
45
Connecting to cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud)|198.23.119.245|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 36689578 (35M) [application/zip]
Saving to: 'employeesdb.zip'

employeesdb.zip      100%[=====] 34.99M  30.3MB/s   in 1.2s

2021-10-12 20:08:25 (30.3 MB/s) - 'employeesdb.zip' saved [36689578/36689578]

theia@theiadocker-:/home/project$
```

2. Next, we'll need to unzip its contents. We can do that with the following command:

```
1. unzip employeesdb.zip
```

```
theia@theiadocker-:/home/project$ unzip employeesdb.zip
Archive: employeesdb.zip
  creating: employeesdb/
  creating: employeesdb/sakila/
  inflating: employeesdb/load_salaries2.dump
  inflating: employeesdb/test_versions.sh
  inflating: employeesdb/objects.sql
  inflating: employeesdb/load_salaries3.dump
  inflating: employeesdb/load_dept_emp.dump
  inflating: employeesdb/test_employees_sha.sql
  inflating: employeesdb/change_log
  creating: employeesdb/images/
  inflating: employeesdb/employees_partitioned_5.1.sql
  inflating: employeesdb/test_employees_md5.sql
  inflating: employeesdb/README.md
  inflating: employeesdb/employees.sql
  inflating: employeesdb/load_titles.dump
  inflating: employeesdb/employees_partitioned.sql
  inflating: employeesdb/load_dept_manager.dump
  inflating: employeesdb/sql_test.sh
  inflating: employeesdb/load_departments.dump
  inflating: employeesdb/load_salaries1.dump
  inflating: employeesdb/show_elapsed.sql
  inflating: employeesdb/load_employees.dump
  inflating: employeesdb/sakila/README.md
  inflating: employeesdb/sakila/sakila-mv-data.sql
  inflating: employeesdb/sakila/sakila-mv-schema.sql
  inflating: employeesdb/images/employees.jpg
  inflating: employeesdb/images/employees.png
  inflating: employeesdb/images/employees.gif
theia@theiadocker-:/home/project$
```

3. Now, let's change directories so that we're able to access the files in the newly created `employeesdb` folder.

```
1. cd employeesdb
```

```
theia@theiadocker-:/home/project$ cd employeesdb
theia@theiadocker-:/home/project$ unzip employeesdb.zip
Archive: employeesdb.zip
  creating: employeesdb/
  creating: employeesdb/sakila/
  inflating: employeesdb/load_salaries2.dump
  inflating: employeesdb/test_versions.sh
  inflating: employeesdb/objects.sql
  inflating: employeesdb/load_salaries3.dump
  inflating: employeesdb/load_dept_emp.dump
  inflating: employeesdb/test_employees_sha.sql
  inflating: employeesdb/change_log
  creating: employeesdb/images/
  inflating: employeesdb/employees_partitioned_5.1.sql
  inflating: employeesdb/test_employees_md5.sql
  inflating: employeesdb/README.md
  inflating: employeesdb/employees.sql
  inflating: employeesdb/load_titles.dump
  inflating: employeesdb/employees_partitioned.sql
  inflating: employeesdb/load_dept_manager.dump
  inflating: employeesdb/sql_test.sh
  inflating: employeesdb/load_departments.dump
  inflating: employeesdb/load_salaries1.dump
  inflating: employeesdb/show_elapsed.sql
  inflating: employeesdb/load_employees.dump
  inflating: employeesdb/sakila/README.md
  inflating: employeesdb/sakila/sakila-mv-data.sql
  inflating: employeesdb/sakila/sakila-mv-schema.sql
  inflating: employeesdb/images/employees.jpg
  inflating: employeesdb/images/employees.png
  inflating: employeesdb/images/employees.gif
theia@theiadocker-:/home/project$ cd employeesdb
theia@theiadocker-:/home/project/employeesdb$
```

4. In order to import the data, we'll need to load the data through MySQL. We can do that by navigating to the `Skills Network Toolbars`, selecting `Databases` and then selecting `MySQL`.

Press **Start**. This will start a session of MySQL in SN Labs.


```
mysql> show tables;
+-----+
| Tables_in_employees |
+-----+
| current_dept_emp      |
| departments           |
| emp_emp               |
| emp_emp_latest_date   |
| emp_manager           |
| employees             |
| job_history            |
| titles                |
+-----+

8 rows in set (0.01 sec)
```

In this database, there are 8 tables, which we can confirm with the database's ERD.

Now that your database is all set up, let's take a look at how we can check a query's performance!

Exercise 2: Check Your Query's Performance with EXPLAIN

The `EXPLAIN` statement, which provides information about how MySQL executes your statement, will offer you insight about the number of rows your query is planning on looking through. This statement can be helpful when your query is running slow. For example, is it running slow because it's scanning the entire table each time?

1. Let's start with selecting all the data from the `employees` table:

```
1. 1
2. SELECT * FROM employees;
Copied!

mysql> head -n 1000000000 /home/mysqltest/db-master <
+-----+-----+-----+-----+-----+-----+
| emp_no | birth_date | first_name | last_name | gender | hire_date |
+-----+-----+-----+-----+-----+-----+
| 499555 | 1961-10-23 | Kananahalli | Zucker    | M      | 1993-03-06 |
| 499556 | 1958-05-06 | Panagosa   | Crooks    | F      | 1994-10-12 |
| 499557 | 1954-10-08 | Simon      | Keshava   | M      | 1987-02-23 |
| 499558 | 1957-05-08 | Strachan   | Theoretbacher | M      | 1993-12-17 |
| 499559 | 1954-01-29 | Lillian    | Setu      | M      | 1991-11-06 |
| 499560 | 1952-12-02 | Graham    | Vachek    | F      | 1989-09-16 |
| 499561 | 1962-02-28 | Moberg    | Hurd      | F      | 1993-07-24 |
| 499562 | 1954-03-28 | Yongplata | Dalton    | M      | 1995-06-28 |
| 499563 | 1954-01-24 | Samy      | Lematt    | M      | 1995-06-29 |
| 499564 | 1958-02-24 | Randy      | Matrov    | M      | 1988-11-18 |
| 499565 | 1961-12-07 | Rompage    | Murray    | F      | 1993-06-16 |
| 499566 | 1955-12-04 | Minelli    | Crabtree  | F      | 1985-06-13 |
| 499567 | 1954-06-23 | Banziling  | Bonoff    | M      | 1986-06-15 |
| 499568 | 1953-03-07 | Dharmaraja | Eril      | M      | 1981-10-06 |
| 499569 | 1948-06-02 | Maxena     | Ducloy    | M      | 1992-02-16 |
| 499570 | 1964-05-25 | Jemali     | Hecavat   | M      | 1988-08-06 |
| 499571 | 1982-12-28 | Uwe        | Ullong    | M      | 1989-02-24 |
| 499572 | 1957-07-25 | Karpas     | Loucky    | M      | 1991-11-11 |
| 499573 | 1961-06-03 | Label      | Tubbano   | M      | 1984-02-01 |
| 499574 | 1954-06-18 | Shuchita   | Piazza    | F      | 1989-09-16 |
| 499575 | 1952-11-09 | Maxwell    | Clervant  | M      | 1992-03-13 |
| 499576 | 1954-06-28 | Gurnahang | Felner    | M      | 1981-10-06 |
| 499577 | 1954-06-04 | Marital    | Kuzuka    | M      | 1989-09-24 |
| 499578 | 1948-03-29 | Chiranjit  | Kuzushiki | M      | 1989-09-24 |
| 499579 | 1954-02-28 | Pradeepam  | Kuzushiki | M      | 1989-09-24 |
| 499580 | 1954-06-26 | Gino       | Hecavat   | M      | 1991-02-11 |
| 499581 | 1951-01-02 | Yumika     | Kishita   | F      | 1993-03-07 |
| 499582 | 1954-06-25 | Mohamed    | Pissukun  | M      | 1986-02-15 |
| 499583 | 1954-06-28 | Uzi        | Jureja    | F      | 1989-08-16 |
| 499584 | 1954-08-31 | Kellogg    | Ruman     | M      | 1985-09-11 |
| 499585 | 1954-12-28 | Gita       | Lukewermeir | M      | 1992-02-11 |
| 499586 | 1951-07-22 | Nathan     | Bana      | F      | 1985-02-11 |
| 499587 | 1961-09-05 | Rini      | Durkin    | F      | 1989-09-28 |
| 499588 | 1962-09-21 | Banziling  | Klerker    | F      | 1986-06-06 |
| 499589 | 1954-06-06 | Kelchiro   | Litchavit  | M      | 1991-10-18 |
| 499590 | 1951-11-03 | Huiang     | Kooling    | M      | 1991-10-18 |
| 499591 | 1954-02-28 | Poonam     | Poonam     | F      | 1987-05-18 |
| 499592 | 1968-10-12 | Simona     | Salverda   | F      | 1987-05-18 |
| 499593 | 1951-06-04 | DeForest   | Mullisathen | M      | 1990-04-24 |
| 499594 | 1952-02-26 | Maxin      | Litcher    | F      | 1993-01-22 |
| 499595 | 1953-03-07 | Zita       | Baaz      | M      | 1989-09-27 |
| 499596 | 1951-06-03 | Barbara    | Lematt    | M      | 1986-04-21 |
| 499597 | 1954-08-05 | Patricia   | Bruppi    | M      | 1993-10-13 |
| 499598 | 1954-05-01 | Sachin     | Tachuda   | M      | 1997-11-30 |
| 499599 | 1954-05-01 | Sachin     | Tachuda   | M      | 1997-11-30 |
+-----+-----+-----+-----+-----+-----+

499600 rows in set (0.34 sec)
```

As you can see, all 500,024 rows were loaded, taking about 0.34 seconds.

2. We can use `EXPLAIN` to see how many rows were scanned:

```
1. 1
2. EXPLAIN SELECT * FROM employees;
Copied!

mysql> head -n 1000000000 /home/mysqltest/db-master <
+-----+-----+-----+-----+-----+-----+
| emp_no | birth_date | first_name | last_name | gender | hire_date |
+-----+-----+-----+-----+-----+-----+
| 499565 | 1961-10-23 | Kananahalli | Zucker    | M      | 1993-03-06 |
| 499566 | 1958-05-06 | Panagosa   | Crooks    | F      | 1994-10-12 |
| 499567 | 1954-10-08 | Simon      | Keshava   | M      | 1987-02-23 |
| 499568 | 1957-05-08 | Strachan   | Theoretbacher | M      | 1993-12-17 |
| 499569 | 1954-01-29 | Lillian    | Setu      | M      | 1991-11-06 |
| 499570 | 1952-12-02 | Graham    | Vachek    | F      | 1989-09-16 |
| 499571 | 1962-02-28 | Moberg    | Hurd      | F      | 1993-07-24 |
| 499572 | 1954-03-28 | Yongplata | Dalton    | M      | 1995-06-28 |
| 499573 | 1954-01-24 | Samy      | Lematt    | M      | 1995-06-29 |
| 499574 | 1958-02-24 | Randy      | Matrov    | M      | 1988-11-18 |
| 499575 | 1961-12-07 | Rompage    | Murray    | F      | 1993-06-16 |
| 499576 | 1955-12-04 | Minelli    | Crabtree  | F      | 1985-06-13 |
| 499577 | 1954-06-23 | Banziling  | Bonoff    | M      | 1986-06-15 |
| 499578 | 1953-03-07 | Dharmaraja | Eril      | M      | 1981-10-06 |
| 499579 | 1948-06-02 | Maxena     | Ducloy    | M      | 1992-02-16 |
| 499580 | 1964-05-25 | Jemali     | Hecavat   | M      | 1988-08-06 |
| 499581 | 1982-12-28 | Uwe        | Ullong    | M      | 1989-02-24 |
| 499582 | 1957-07-25 | Karpas     | Loucky    | M      | 1991-11-11 |
| 499583 | 1961-06-03 | Label      | Tubbano   | M      | 1984-02-01 |
| 499584 | 1954-06-18 | Shuchita   | Piazza    | F      | 1989-09-16 |
| 499585 | 1952-11-09 | Maxwell    | Clervant  | M      | 1992-03-13 |
| 499586 | 1954-06-28 | Gurnahang | Felner    | M      | 1981-10-06 |
| 499587 | 1954-06-04 | Marital    | Kuzuka    | M      | 1989-09-24 |
| 499588 | 1948-03-29 | Chiranjit  | Kuzushiki | M      | 1989-09-24 |
| 499589 | 1954-02-28 | Pradeepam  | Kuzushiki | M      | 1989-09-24 |
| 499590 | 1954-06-26 | Gino       | Hecavat   | M      | 1991-02-11 |
| 499591 | 1951-01-02 | Yumika     | Kishita   | F      | 1993-03-07 |
| 499592 | 1954-06-25 | Mohamed    | Pissukun  | M      | 1986-02-15 |
| 499593 | 1954-06-28 | Uzi        | Jureja    | F      | 1989-08-16 |
| 499594 | 1954-08-31 | Kellogg    | Ruman     | M      | 1985-09-11 |
| 499595 | 1954-12-28 | Gita       | Lukewermeir | M      | 1992-02-11 |
| 499596 | 1951-07-22 | Nathan     | Bana      | F      | 1985-02-11 |
| 499597 | 1961-09-05 | Rini      | Durkin    | F      | 1989-09-28 |
| 499598 | 1962-09-21 | Banziling  | Klerker    | F      | 1986-06-06 |
| 499599 | 1954-06-06 | Kelchiro   | Litchavit  | M      | 1991-10-18 |
| 499600 | 1951-11-03 | Huiang     | Kooling    | M      | 1991-10-18 |
| 499601 | 1954-02-28 | Poonam     | Poonam     | F      | 1987-05-18 |
| 499602 | 1968-10-12 | Simona     | Salverda   | F      | 1987-05-18 |
| 499603 | 1951-06-04 | DeForest   | Mullisathen | M      | 1990-04-24 |
| 499604 | 1952-02-26 | Maxin      | Litcher    | F      | 1993-01-22 |
| 499605 | 1953-03-07 | Zita       | Baaz      | M      | 1989-09-27 |
| 499606 | 1951-06-03 | Barbara    | Lematt    | M      | 1986-04-21 |
| 499607 | 1954-08-05 | Patricia   | Bruppi    | M      | 1993-10-13 |
| 499608 | 1954-05-01 | Sachin     | Tachuda   | M      | 1997-11-30 |
| 499609 | 1954-05-01 | Sachin     | Tachuda   | M      | 1997-11-30 |
+-----+-----+-----+-----+-----+-----+

499610 rows in set (0.37 sec)
```

```
mysql> EXPLAIN SELECT * FROM employees;
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1  | SIMPLE      | employees | NULL | ALL | NULL | NULL | NULL | NULL | 299423 | 33.33 | Using where |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

1 row in set, 1 warning (0.00 sec)
```

Notice how `rows_examined` shows that it is examining 299,980 rows, almost the entire table! With a larger table, this could result in the query running slowly.

So, how can we make this query faster? That's where indexes come in!

Exercise 3: Add an Index to Your Table

1. To begin, let's take at the existing indexes. We can do that by entering the following command:

```
1. 1
2. SHOW INDEX FROM employees;
Copied!

mysql> mysql> show indexes from employees;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| Table | Non_unique | Key_name | Seq_in_index | Column_name | Collation | Cardinality | Sub_part | Packed | Null | Index_type | Comment | Index_comment | Visible | Expression |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| employees | 0 | PRIMARY | 1 | emp_no | A | 299423 | NULL | NULL | NULL | BTREE | | | YES | NULL |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

1 row in set (0.00 sec)
```

```
mysql>
```

Remember that indexes for primary keys are created automatically, as we can see above. An index has already been created for the primary key, `emp_no`. If we think about this, this makes sense because each employee number is unique to the employee, with no NULL values.

2. Now, let's say we wanted to see all the information about employees who were hired on or after January 1, 2000. We can do that with the query:

```
1. 1
2. SELECT * FROM employees WHERE hire_date >= '2000-01-01';
Copied!

mysql> SELECT * FROM employees WHERE hire_date >= '2000-01-01';
+-----+-----+-----+-----+-----+-----+
| emp_no | birth_date | first_name | last_name | gender | hire_date |
+-----+-----+-----+-----+-----+-----+
| 47291 | 1960-09-09 | Ulf        | Flexer    | M      | 2000-01-12 |
| 60134 | 1964-04-21 | Seshu     | Rathonyi  | F      | 2000-01-02 |
| 72329 | 1953-02-09 | Randi     | Luit      | F      | 2000-01-02 |
| 108201 | 1955-04-14 | Mariangiola | Boreale   | M      | 2000-01-01 |
| 205048 | 1960-09-12 | Ennio     | Alblas    | F      | 2000-01-06 |
| 222965 | 1959-08-07 | Volkmar   | Perko     | F      | 2000-01-13 |
| 226633 | 1958-06-10 | Xuejun    | Benzmueller | F      | 2000-01-04 |
| 227544 | 1954-11-17 | Shahab    | Demeyer   | M      | 2000-01-08 |
| 422990 | 1953-04-09 | Jaana     | Verspoor  | F      | 2000-01-11 |
| 424445 | 1953-04-27 | Jeong     | Boreale   | M      | 2000-01-03 |
| 428377 | 1957-05-09 | Yucai     | Gerlach   | M      | 2000-01-23 |
| 463807 | 1964-06-12 | Bikash    | Covnot    | M      | 2000-01-28 |
| 499553 | 1954-05-06 | Hideyuki  | Delgrande | F      | 2000-01-22 |
+-----+-----+-----+-----+-----+-----+

13 rows in set (0.17 sec)
```

As we can see, the 13 rows returned took about 0.17 seconds to execute. That may not seem like a long time with this table, but keep in mind that with larger tables, this time can vary greatly.

3. With the `EXPLAIN` statement, we can check how many rows this query is scanning:

```
1. 1
2. EXPLAIN SELECT * FROM employees WHERE hire_date >= '2000-01-01';
Copied!

mysql> EXPLAIN SELECT * FROM employees WHERE hire_date >= '2000-01-01';
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1  | SIMPLE      | employees | NULL | ALL | NULL | NULL | NULL | NULL | 299423 | 33.33 | Using where |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

1 row in set, 1 warning (0.01 sec)
```

```
mysql>
```

This query results in a scan of 299,423 rows, which is nearly the entire table!

By adding an index to the `hire_date` column, we'll be able to reduce the query's need to search through every entry of the table, instead only searching through what it needs.

4. You can add an index with the following:

```
1. 1
2. CREATE INDEX hire_date_index ON employees(hire_date);
Copied!
```

The `CREATE INDEX` command creates an index called `hire_date_index` on the table `employees` on column `hire_date`.


```
mysql> CREATE INDEX first_name_index ON employees(first_name);
Query OK, 0 rows affected (1.59 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> CREATE INDEX last_name_index ON employees(last_name);
Query OK, 0 rows affected (1.75 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> SHOW INDEX from employees;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment	Visible	Expression
employees	0	PRIMARY	1	emp_no	A	299423				BTREE			YES	NULL
employees	1	first_name_index	1	first_name	A	1251		NULL	NULL	BTREE			YES	NULL
employees	1	last_name_index	1	last_name	A	1585		NULL	NULL	BTREE			YES	NULL

3 rows in set (0.01 sec)

4. Great! With your indexes now in place, we can re-run the query:

```
1. 1
2. SELECT * FROM employees WHERE first_name LIKE 'C%' OR last_name LIKE 'C%';
Copied
```

499881	1952-12-01	Christoph	Schneeberger	F	1987-10-29
499889	1956-01-29	Charlene	Hasham	F	1988-03-19
499908	1953-07-19	Cherone	Coorg	F	1988-12-02
499916	1962-01-09	Florina	Cusworth	F	1997-05-18
499920	1953-07-18	Christ	Murtagh	M	1986-04-17
499933	1957-10-21	Chianti	Riesenhuber	F	1993-05-28
499936	1954-02-11	Chiranjit	Himler	M	1994-10-31
499947	1960-02-06	Conrado	Koyama	F	1989-02-19
499948	1953-05-24	Cordelia	Paludetto	M	1993-01-28
499956	1959-01-08	Zhonghua	Crooks	F	1994-10-12
499966	1955-12-04	Mihailis	Crabtree	F	1985-06-13
499975	1952-11-09	Masali	Chorvat	M	1992-01-23
499978	1960-03-29	Chiranjit	Kuzuoka	M	1990-05-24

28970 rows in set (0.16 sec)

Let's also see how many rows are being scanned:

```
1. 1
2. EXPLAIN SELECT * FROM employees WHERE first_name LIKE 'C%' OR last_name LIKE 'C%';
Copied
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	employees	NULL	ALL	first_name_index, last_name_index	NULL	NULL	NULL	299423	20.99	Using where

1 row in set, 1 warning (0.00 sec)

With indexes, the query still scans all the rows.

5. Let's use the `INDEX` all clause to improve the performance of this query.

We can do this with the following:

```
1. 1
2. SELECT * FROM employees WHERE first_name LIKE 'C%' UNION ALL SELECT * FROM employees WHERE last_name LIKE 'C%';
Copied
```

492481	1953-01-16	Chikara	Czap	M	1990-05-23
496858	1957-12-26	Cheong	Czap	F	1994-10-26

29730 rows in set (0.11 sec)

As we can see, this query only takes 0.11 seconds to execute, running faster than when we used the `OR` operator.

Using the `EXPLAIN` statement, we can see why that might be:

```
mysql> EXPLAIN SELECT * FROM employees WHERE first_name LIKE 'C%' UNION ALL SELECT * FROM employees WHERE last_name LIKE 'C%';
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	PRIMARY	employees	NULL	range	first_name_index	first_name_index	58	NULL	20622	100.00	Using index condition
2	UNION	employees	NULL	range	last_name_index	last_name_index	66	NULL	34168	100.00	Using index condition

2 rows in set, 1 warning (0.00 sec)

As the `EXPLAIN` statement reveals, there were two `SELECT` operations performed, with the total number of rows scanned sitting at 54,790. This is less than the original query that scanned the entire table and, as a result, the query performs faster.

Please note, if you choose to perform a leading wildcard search with an index, the entire table will still be scanned. You can see this yourself with the following query:

```
1. 1
2. SELECT * FROM employees WHERE first_name LIKE 'C%';
Copied
```

With this query, we want to find all the employees whose first names end with "C".

When checking with the `EXPLAIN` and `SHOW INDEX` statements, we can see that although we have an index on `first_name`, the index is not used and results in a search of the entire table.

Under the `EXPLAIN` statement's `possible_keys` column, we can see that this index has not been used as the entry is `NULL`.

```
mysql> EXPLAIN SELECT * FROM employees WHERE first_name LIKE 'C%';
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	employees	NULL	ALL	NULL	NULL	NULL	NULL	299423	11.11	Using where

1 row in set, 1 warning (0.00 sec)

```
mysql> SHOW INDEX from employees;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment	Visible	Expression
employees	0	PRIMARY	1	emp_no	A	299423				BTREE			YES	NULL
employees	1	first_name_index	1	first_name	A	1251		NULL	NULL	BTREE			YES	NULL
employees	1	last_name_index	1	last_name	A	1585		NULL	NULL	BTREE			YES	NULL

3 rows in set (0.00 sec)

On the other hand, indexes do work with trailing wildcards, as seen with the following query that finds all employees whose first names begin with "C".

```
1. 1
2. SELECT * FROM employees WHERE first_name LIKE 'C%';
Copied
```

492080	1961-08-02	Cullen	Whittlesey	F	1997-01-12
495632	1958-05-16	Cullen	Pollock	M	1992-01-21

11294 rows in set (0.04 sec)

```
mysql> EXPLAIN SELECT * FROM employees WHERE first_name LIKE 'C%';
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	employees	NULL	range	first_name_index	first_name_index	58	NULL	20622	100.00	Using index condition

1 row in set, 1 warning (0.01 sec)

Under the `EXPLAIN` statement's `possible_keys` and `Extra` columns, we can see that the `first_name_index` is used. With only 20,622 rows scanned, the query performs better.

Exercise 5: Be SELECTIVE

In general, it's best practice to only select the columns that you need. For example, if you wanted to see the names and hire dates of the various employees, you could show that with the following query:

```
1. 1
2. SELECT * FROM employees;
Copied
```

499998	1956-09-05	Patricia	Breugel	M	1993-10-13
499999	1958-05-01	Sachin	Tsukuda	M	1997-11-30

300024 rows in set (0.26 sec)

```
mysql> EXPLAIN SELECT * FROM employees;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	employees	NULL	ALL	NULL	NULL	NULL	NULL	299423	100.00	NULL

1 row in set, 1 warning (0.01 sec)

Notice how the query loads 300,024 rows in about 0.26 seconds. With the `EXPLAIN` statement, we can see that the entire table is being scanned, which makes sense because we are looking at all the entries.

If we, however, only wanted to see the names and hire dates, then we should select those columns:

```
1. 1
2. SELECT first_name, last_name, hire_date FROM employees;
```

[Copy](#)

```
| Patricia | Breugel | 1993-10-13 |
| Sachin | Tsukuda | 1997-11-30 |
-----+-----+-----+
300024 rows in set (0.17 sec)

mysql> EXPLAIN SELECT first_name, last_name, hire_date FROM employees;
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | employees | NULL | ALL | NULL | NULL | NULL | NULL | 299423 | 100.00 | NULL |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```

As you can see, this query was executed a little faster despite scanning the entire table as well.

Give this a try!

Practice Exercise 1

Let's take a look at the **salaries** table. What if we wanted to see how much each employee earns?

When running the query, keep in mind how long it takes the query to run and how many rows are scanned each time.

1. First, let's select all the rows and columns from this table.

- [Hint \(Click Here\)](#)
- [Solution \(Click Here\)](#)

To select all the rows and columns, we'll use the following query:

```
1. 1
2. SELECT * FROM salaries;
```

[Copy](#)

Although the exact time may differ, in this instance, it took about 1.71 seconds to load 2,844,047 rows.

We can check how many rows were scanned with the following statement:

```
1. 1
2. EXPLAIN SELECT * FROM salaries;
```

[Copy](#)

We can see that almost the entire table was scanned, as expected, totalling to 2,838,426 rows.

```
| 400000 | 61843 | 1998-11-28 | 1998-11-28 |
| 400000 | 78145 | 1998-11-28 | 2000-11-29 |
| 400000 | 74227 | 2000-11-29 | 2001-11-29 |
| 400000 | 77383 | 2001-11-29 | 2002-01-01 |
-----+-----+-----+-----+
2844047 rows in set (1.71 sec)

mysql> EXPLAIN SELECT * FROM salaries;
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | salaries | NULL | ALL | NULL | NULL | NULL | NULL | 2838426 | 100.00 | NULL |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```

2. Now, let's see if there's a way to optimize this query. Since we only want to see how much each employee earns, then we can just select a few columns instead of all of them. Which ones would you select?

- [Hint \(Click Here\)](#)
- [Solution \(Click Here\)](#)

To select columns that will give us information about the employee and their corresponding salary, we'll choose the **emp_no** and **salary** columns with the following query:

```
1. 2
2. SELECT emp_no, salary FROM salaries;
```

[Copy](#)

Although the exact time may differ, in this instance, it took about 1.19 seconds to load 2,844,047 rows.

We can check how many rows were scanned with the following statement:

```
1. 2
2. EXPLAIN SELECT emp_no, salary FROM salaries;
```

[Copy](#)

We can see that almost the entire table was scanned, as expected, totalling to 2,838,426 rows. Yet, it loaded faster than the first instance because we were more selective in the columns that were chosen.

```
| 400000 | 61843 |
| 400000 | 78145 |
| 400000 | 74227 |
| 400000 | 77383 |
-----+-----+
2844047 rows in set (1.19 sec)

mysql> EXPLAIN SELECT emp_no, salary FROM salaries;
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | salaries | NULL | ALL | NULL | NULL | NULL | NULL | 2838426 | 100.00 | NULL |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```

Practice Exercise 2

Let's take a look at the **titles** table. What if we wanted to see the employee and their corresponding title?

Practice by selecting only the necessary columns and run the query!

- [Hint \(Click Here\)](#)
- [Solution \(Click Here\)](#)

To select columns that will give us information about the employee and their corresponding title, we'll choose the **emp_no** and **title** columns with the following query:

```
1. 2
2. SELECT emp_no, title FROM titles;
```

[Copy](#)

Although the exact time may differ, in this instance, it took about 0.22 seconds to load 443,308 rows.

We can check how many rows were scanned with the following statement:

```
1. 2
2. EXPLAIN SELECT emp_no, title FROM titles;
```

[Copy](#)

We can see that almost the entire table was scanned, as expected, totalling to 442,545 rows.

```
| 400001 | Senior Engineer |
| 400008 | Senior Staff |
| 400009 | Staff |
| 400019 | Engineer |
-----+-----+
443308 rows in set (0.22 sec)

mysql> EXPLAIN SELECT emp_no, title FROM titles;
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | titles | NULL | index | NULL | PRIMARY | 200 | NULL | 442545 | 100.00 | Using index |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```

In comparison, if you had run this with all columns selected, you may have noticed that it took about 0.47 seconds to load and scan the same amount of rows.

```
| 400001 | Senior Engineer | 1997-08-29 | 9999-01-01 |
| 400008 | Senior Staff | 1996-12-27 | 9999-01-01 |
| 400009 | Staff | 1995-12-27 | 1996-11-27 |
| 400019 | Engineer | 1997-11-28 | 9999-01-01 |
-----+-----+-----+-----+
443308 rows in set (0.47 sec)

mysql> EXPLAIN SELECT * FROM titles;
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | titles | NULL | ALL | NULL | NULL | NULL | NULL | 442545 | 100.00 | NULL |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```

Conclusion

Congratulations! Now, not only can you now identify common causes to slow queries, but you can resolve them by applying the knowledge that you have gained in this lab. Equipped with this problem-solving skill, you will be able to improve your queries performance, even in large databases.

Author(s)

Kathy An

Other Contributor(s)

Ravi Ahuja

Changelog

Date	Version	Changed by	Change Description
2021-10-05 1.0		Kathy An	Created initial version
2022-09-08 1.1		Lakshmi Hella	Made changes in practice exercise
2023-05-08 1.2		Evin Hao	Updated Page Frames

© IBM Corporation 2023. All rights reserved.